


Why Repeated Judges Matter

One open-ended judgment per condition cannot tell apart a prompt effect from random judge variation.
Repeated, independent judgments can.

One judgment per condition

You only see one voice.

Was the difference real, or just this judge?



We can't tell:


- The prompt changed the result.
- Or this judge just had an unusual opinion.

→ Attribution is impossible.

Repeated, independent judgments


You hear many independent voices.

Real effects show up. Random noise averages out.



Compare the distributions, not single voices.

- If they are clearly different → likely a prompt effect.
- If they overlap heavily → likely random variation.



Takeaway

One judgment is a coin flip. 
Many independent judgments reveal the signal.



Goal: separate
signal from noise.



Better design.
Stronger claims.